

Operating System Lab-8

Q1> Consider a system with one parent process and two child processes A and B. There is a shared signed integer X initialized to 0. Process A increments X by one, 100000 times in a for loop. Process B decrements X by one, 100000 times in a for loop. After both A and B finish, the parent process prints the final value of X.

Hint: Declare a shared memory variable to hold X (using system calls `shmget()`, `shmat()`, `shmdt()`, and `shmctl()`). Write the programs for processes A and B. Do not put any synchronization code in the code for A and B. You should write the code in such a way so that you can simulate race condition in your program by slowing down A or B appropriately by using `sleep()` calls at appropriate points. Note that if there is no race condition, the value of X finally should be 0. Simulating race conditions means that if you run the program a few times, sometimes the final value of X printed by your program should be non-zero.

Desired Output:

```
vboxuser@blockchainexplorer:~$ ./race
Final value of X: 0
vboxuser@blockchainexplorer:~$ ./race
Final value of X: -703
vboxuser@blockchainexplorer:~$ ./race
Final value of X: -8
vboxuser@blockchainexplorer:~$ ./race
Final value of X: -3
vboxuser@blockchainexplorer:~$ ./race
Final value of X: 7
vboxuser@blockchainexplorer:~$ █
```

Q2> Add synchronization code based on semaphores to process A and B in question 1 so that there is no possibility of race conditions. Use the calls `semget()`, `semop()`, `semctl()` to create and manage semaphores.

Desired Output:

```
vboxuser@blockchainexplorer:~$ ./race2
Final value of X: 0
vboxuser@blockchainexplorer:~$ ./race2
Final value of X: 0
vboxuser@blockchainexplorer:~$ ./race2
Final value of X: 0
vboxuser@blockchainexplorer:~$ ./race2
Final value of X: 0
vboxuser@blockchainexplorer:~$ ./race2
Final value of X: 0
vboxuser@blockchainexplorer:~$
```

Note:

The important headers to be included are:

- **<sys/types.h>:** Contains various data types and symbolic constants used in system calls. This header file is often included to ensure portability of code between different Unix-like systems.
- **<sys/ipc.h>:** Provides definitions for interprocess communication data structures. It includes structures and constants used with functions like `ftok()` for generating unique keys for IPC objects.
- **<sys/shm.h>:** Contains definitions for shared memory operations. Shared memory is a way for processes to share data by mapping a common memory area into their address spaces. Functions like `shmget()`, `shmat()`, and `shmdt()` are used for shared memory operations.
- **<sys/sem.h>:** Contains definitions for semaphores, which are synchronization mechanisms used for process coordination. Functions like `semget()`, `semop()`, and `semctl()` are used to work with semaphores.
- **<sys/wait.h>:** Provides functions and constants related to process management, specifically waiting for child processes to terminate. Functions like `wait()`, `waitpid()`, and macros like `WEXITSTATUS` are commonly used for process control and synchronization.

The system calls to be used for performing the above tasks are as follows: -

- **shmget():** `shmget()` is a system call in Unix-like operating systems (including Linux) that is used for working with shared memory segments. Shared memory is a form of inter-process communication (IPC) that allows multiple processes to share a common memory region. `shmget()` is used for creating, accessing, or

manipulating shared memory segments. The syntax for the shmget() function is as follows:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

```
int shmget(key_t key, size_t size, int shmflg);
```

where,

key: This is an identifier for the shared memory segment. It can be generated using the ftok function, which converts a pathname and a project identifier into a key. The same key is used to identify and access the shared memory segment across different processes.

size: This parameter specifies the size (in bytes) of the shared memory segment you want to create or access.

shmflg: This is a set of flags that control the behavior of the shmget function.

- **shmat():** The shmat() function is used in Unix-like operating systems (including Linux) for attaching a shared memory segment to the address space of a process. shmat() is used to establish a connection between a process and a shared memory segment, allowing the process to read from and write to that shared memory. The syntax for shmat() function is as follows:-

```
#include <sys/types.h>
#include <sys/shm.h>
```

```
void *shmat(int shmid, const void *shmaddr, int shmflg);
```

where,

shmid: This is the identifier of the shared memory segment, typically obtained using shmget(). It identifies the specific shared memory segment you want to attach to.

shmaddr: This parameter specifies the desired address at which you want to attach the shared memory segment in the process's address space. If you pass NULL, the system chooses a suitable address for attachment.

shmflg: This is a set of flags that control the behavior of the shmat() function.

The return value of `shmat()` is a pointer to the attached shared memory segment in the process's address space. If the attachment is unsuccessful, `shmat()` returns `(void*) -1`, and you can check the specific error condition by examining the `errno` variable.

- **shmdt():** The `shmdt()` function is used in Unix-like operating systems (including Linux) for detaching a process from a previously attached shared memory segment. When a process is done using the shared memory, it should detach from it using `shmdt()` to release the associated resources. The syntax for `shmdt()` function is as follows:

```
#include <sys/types.h>
#include <sys/shm.h>
```

```
int shmdt(const void *shmaddr);
```

where,

shmaddr: This is a pointer to the address where the shared memory segment is attached in the process's address space. You provide the same pointer that was returned by the `shmat()` function when you originally attached to the shared memory segment. The return value of `shmdt()` is an integer. It returns 0 if the operation is successful (i.e., the process successfully detaches from the shared memory segment). If there's an error, it returns -1, and you can check the specific error condition by examining the `errno` variable.

- **shmctl():** The `shmctl()` function is used in Unix-like operating systems (including Linux) for controlling and managing shared memory segments. Shared memory is a form of inter-process communication (IPC) that allows multiple processes to share a common memory region. `shmctl()` provides a means to perform various operations on shared memory segments, including creating, deleting, and modifying their properties. The syntax for `shmctl` function is as follows:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

```
int shmctl(int shmid, int cmd, struct shmid_ds *buf);
```

where,

shmid: This is the identifier of the shared memory segment you want to perform operations on. You typically obtain this identifier using `shmget()` when creating or accessing a shared memory segment.

cmd: This parameter specifies the command or operation you want to perform on the shared memory segment. The possible values for `cmd` are defined as constants, including:

IPC_STAT: Retrieve information about the shared memory segment (fills `buf` with the information).

IPC_SET: Set information and permissions for the shared memory segment based on the values in `buf`.

IPC_RMID: Remove (delete) the shared memory segment.

buf: This is a pointer to a structure of type `struct shmid_ds` where information about the shared memory segment is stored or from which information is read, depending on the value of `cmd`. The structure may contain data such as segment size, permissions, user-defined data, etc.